

Ensemble method

Geena Kim



Ensemble method- review

Problem: Trees are weak learner and trees overfit

Idea 1: Let's average them (Ensemble)

Idea 2: Let's make decorrelated trees (samples, features)



Random Forest

Random Forest

Bagging : random sampling of data

+

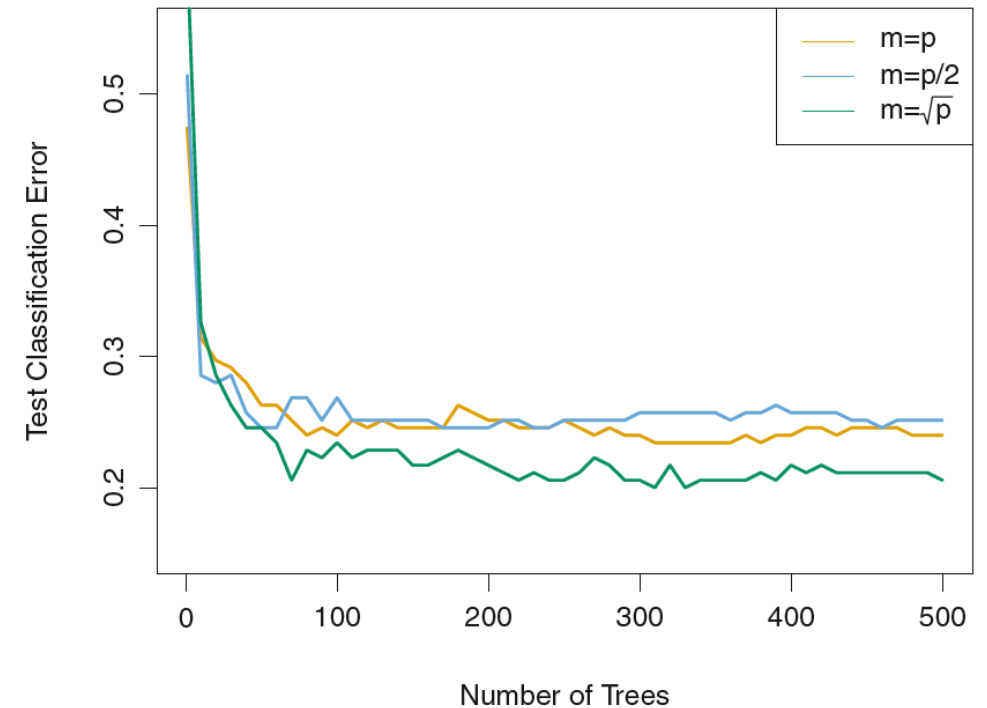
Decorrelation : random sampling of features

II

Random Forest

How do we sample features?

-> Rule of thumb : \sqrt{n}



Ensemble method- review

Problem: Trees are weak learner and trees overfit

Idea 3: Let's make the trees a strong learner

How: Grow a small tree (stump) to fit residual



Boosting

Boosting

1. Set $\hat{f}(x) = 0$ and $r_i = y_i$ for all i in the training set.
2. For $b = 1, 2, \dots, B$ repeat:
 - (a) Fit a tree \hat{f}^b with d splits ($d + 1$ terminal nodes) to the training data (X, r) .
 - (b) Update \hat{f} by adding in a shrunken version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x).$$

- (c) Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i).$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x).$$

shrinkage parameter = hyper
 $0 < \lambda < 1$
 $\rightarrow 0.01$
 0.001

3 HyperParams

λ

B

$d = 1, 2$

Popular Boosted Tree Methods

- AdaBoost (Adaptive Boosting)
- GBM (Gradient Boosting Machine)
- XGBoost (Extreme Gradient Boosting) ←

Popular Boosted Tree Methods

- AdaBoost (Adaptive Boosting)
- GBM (Gradient Boosting Machine)
- XGBoost (Extreme Gradient Boosting)

AdaBoost

Idea: Focus on the misclassified samples

→ Initialize data weights to $w_i = \frac{1}{m}, i = 1, \dots, m$

For $k = 1$ to K :

Fit estimator $f_k(\mathbf{x})$ to training data with weights w_i

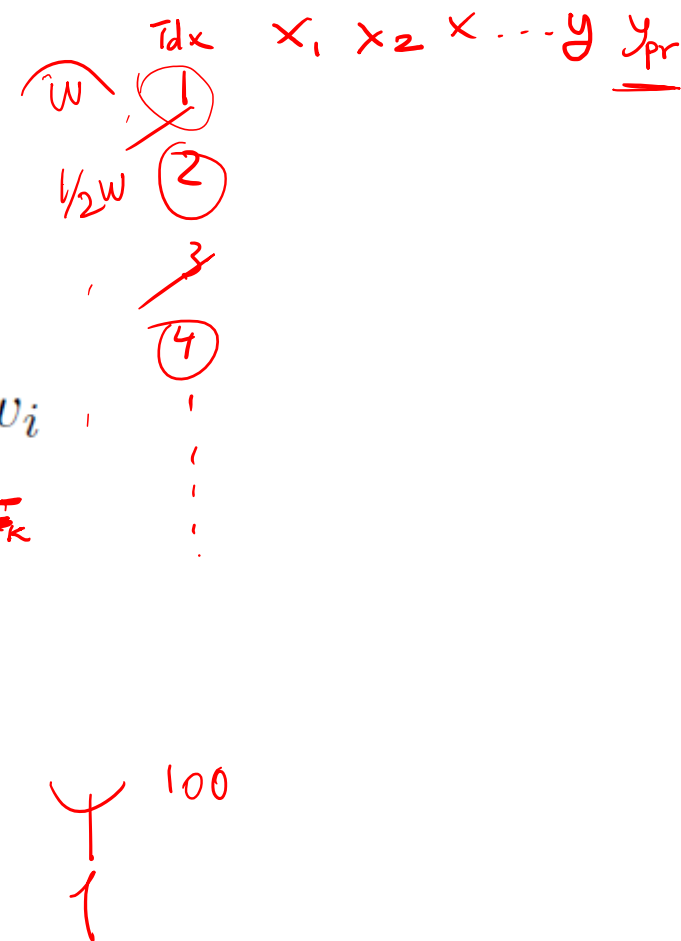
Compute weighted error $\epsilon_k = \frac{\sum_{i=1}^m w_i \mathbb{I}(y_i \neq f_k(\mathbf{x}_i))}{\sum_{i=1}^m w_i}$ $y_{rr} \leftarrow F_k$

Compute estimator weight $\lambda_k = \frac{1}{2} \log((1 - \epsilon_k)/\epsilon_k)$

Update sample weight $w_i \leftarrow \frac{w_i}{Z_k} \cdot \exp[-\lambda_k y_i f_k(\mathbf{x}_i)]$

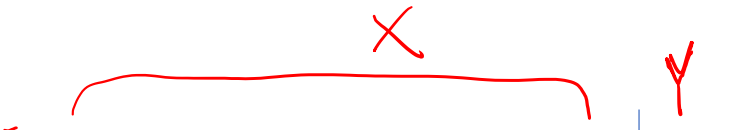
Final model $F(\mathbf{x}) = \text{sign} \left[\sum_{k=1}^K \lambda_k f_k(\mathbf{x}) \right]$

$$F_K = \sum_{k=1}^K \lambda_k f_k$$



AdaBoost

Initialize data weights to $w_i = \frac{1}{m}, i = 1, \dots, m$




<i>i</i>	Age	Sex	ChestPain	Chol	AHD
0	63	1	typical	233	No
1	67	1	asymptomatic	286	Yes
2	67	1	asymptomatic	229	Yes
3	37	1	nonanginal	250	No
4	41	0	nontypical	204	No
5	56	1	nontypical	236	No
6	62	0	asymptomatic	268	Yes
7	57	0	asymptomatic	354	No
8	63	1	asymptomatic	254	Yes
9	53	1	asymptomatic	203	Yes

AdaBoost

$k > 1$

Fit estimator $f_k(\mathbf{x})$ to training data with weights w_i



	Age	Sex	ChestPain	Chol	AHD	weight	Yp
0	63	1	typical	233	No	0.1	Yes
1	67	1	asymptomatic	286	Yes	0.1	Yes
2	67	1	asymptomatic	229	Yes	0.1	Yes
3	37	1	nonanginal	250	No	0.1	No
4	41	0	nontypical	204	No	0.1	No
5	56	1	nontypical	236	No	0.1	No
6	62	0	asymptomatic	268	Yes	0.1	Yes
7	57	0	asymptomatic	354	No	0.1	No
8	63	1	asymptomatic	254	Yes	0.1	Yes
9	53	1	asymptomatic	203	Yes	0.1	No
10	57	1	asymptomatic	192	No	0.1	No

AdaBoost

Compute weighted error $\epsilon_k = \frac{\sum_{i=1}^m w_i \mathbb{I}(y_i \neq f_k(\mathbf{x}_i))}{\sum_{i=1}^m w_i}$

Handwritten notes: Red arrows point from the numerator to the denominator. Above the denominator, '+1, -1' is written. A red circle is drawn around the first 'w_i' in the numerator.

	Age	Sex	ChestPain	Chol	AHD	weight	Yp
0	63	1	typical	233	No	0.1	Yes
1	67	1	asymptomatic	286	Yes	0.1	Yes
2	67	1	asymptomatic	229	Yes	0.1	Yes
3	37	1	nonanginal	250	No	0.1	No
4	41	0	nontypical	204	No	0.1	No
5	56	1	nontypical	236	No	0.1	No
6	62	0	asymptomatic	268	Yes	0.1	Yes
7	57	0	asymptomatic	354	No	0.1	No
8	63	1	asymptomatic	254	Yes	0.1	Yes
9	53	1	asymptomatic	203	Yes	0.1	No
10	57	1	asymptomatic	192	No	0.1	No

$$\epsilon_k = 0.2$$

Handwritten notes: A red '-1' is written next to the first row (index 0). A red '1' is written next to the second row (index 1).

AdaBoost

Compute estimator weight $\lambda_k = \frac{1}{2} \log((1 - \epsilon_k)/\epsilon_k)$

	Age	Sex	ChestPain	Chol	AHD	weight	Yp
0	63	1	typical	233	No	0.1	Yes
1	67	1	asymptomatic	286	Yes	0.1	Yes
2	67	1	asymptomatic	229	Yes	0.1	Yes
3	37	1	nonanginal	250	No	0.1	No
4	41	0	nontypical	204	No	0.1	No
5	56	1	nontypical	236	No	0.1	No
6	62	0	asymptomatic	268	Yes	0.1	Yes
7	57	0	asymptomatic	354	No	0.1	No
8	63	1	asymptomatic	254	Yes	0.1	Yes
9	53	1	asymptomatic	203	Yes	0.1	No
10	57	1	asymptomatic	192	No	0.1	No

$$\lambda_k = 0.69$$

AdaBoost

Update sample weight $w_i \leftarrow \frac{w_i}{Z_k} \exp[\ominus \lambda_k y_i f_k(\mathbf{x}_i)]$

$e^a \rightarrow e^p$

$y \neq f_k \Rightarrow -1$

$y = f_k \Rightarrow 1 \quad e^{-0}$

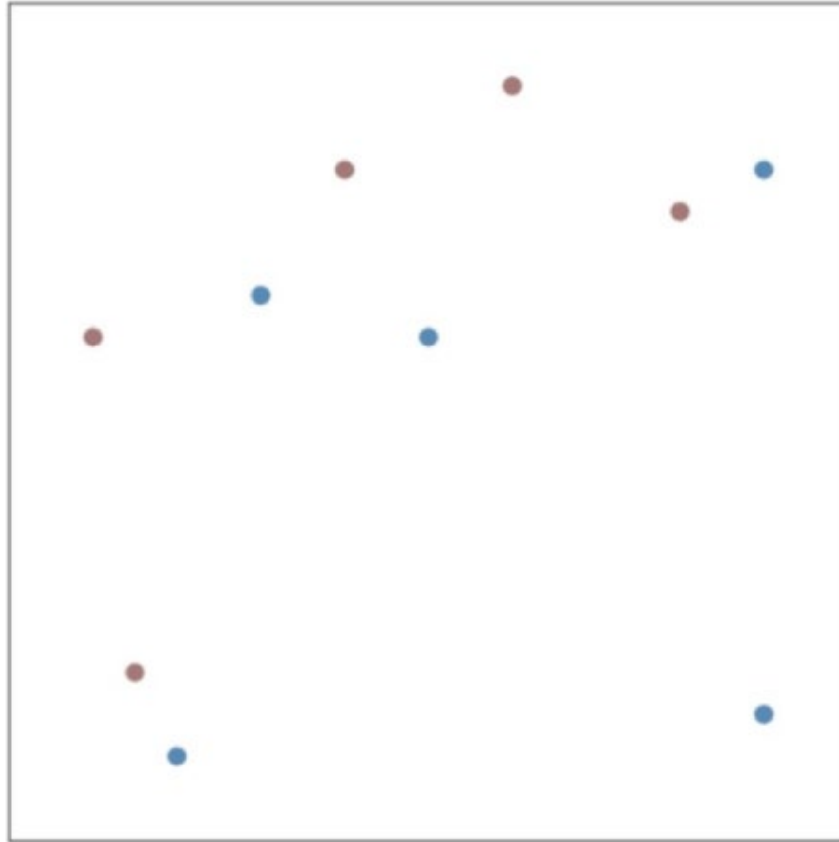
	Age	Sex	ChestPain	Chol	AHD	weight	
✓	0	63	1	typical	233	No	0.2500
-	1	67	1	asymptomatic	286	Yes	0.0625
-	2	67	1	asymptomatic	229	Yes	0.0625
-	3	37	1	nonanginal	250	No	0.0625
	4	41	0	nontypical	204	No	0.0625
	5	56	1	nontypical	236	No	0.0625
	6	62	0	asymptomatic	268	Yes	0.0625
	7	57	0	asymptomatic	354	No	0.0625
	8	63	1	asymptomatic	254	Yes	0.0625
✓	9	53	1	asymptomatic	203	Yes	0.2500
	10	57	1	asymptomatic	192	No	0.1000

Repeat until K or the error =0

Final model $F(\mathbf{x}) = \text{sign} \left[\sum_{k=1}^K \lambda_k f_k(\mathbf{x}) \right]$

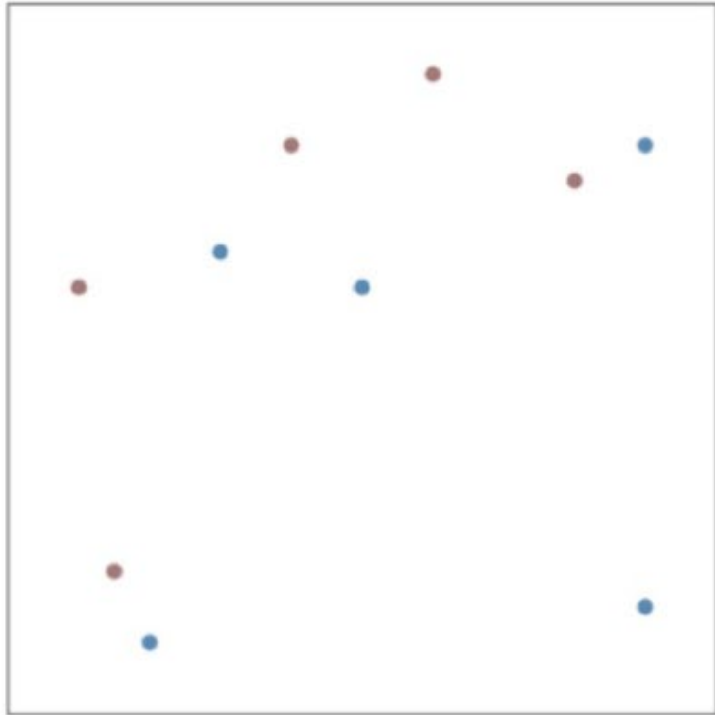
AdaBoost

For a train data

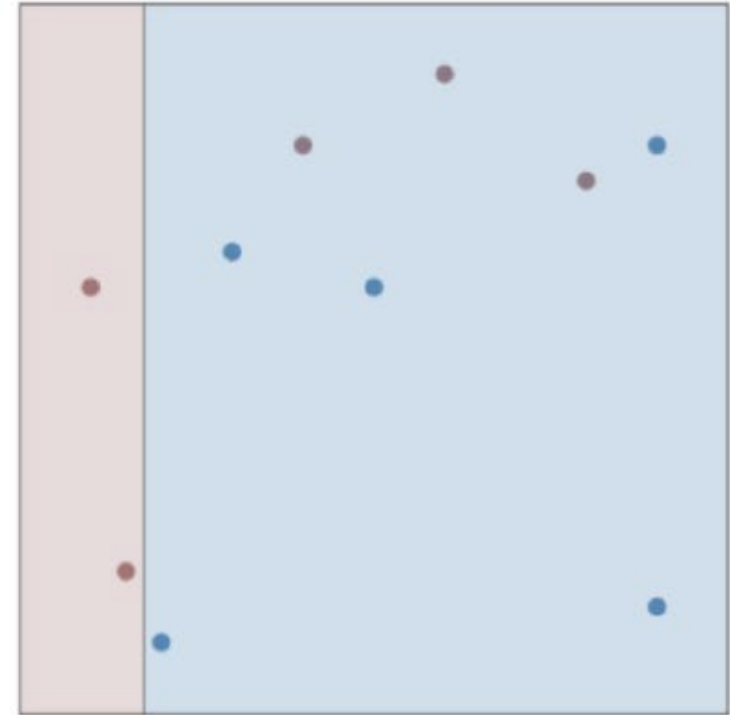


AdaBoost

Fit first stump

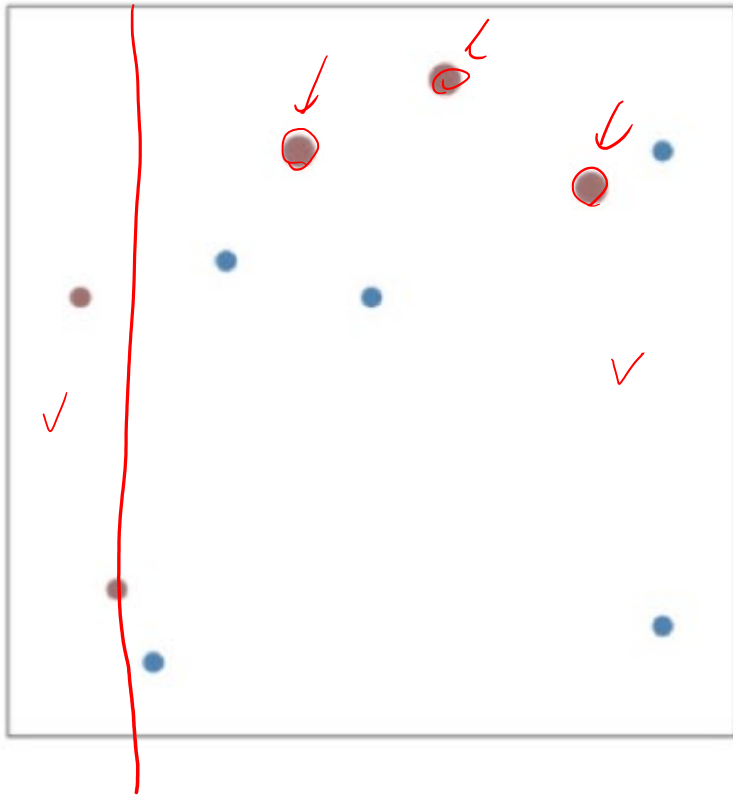


$k = 1, \epsilon = 0.3, \lambda = 0.42$




AdaBoost

Fit second stump




$k = 2, \epsilon = 0.21, \lambda = 0.65$




AdaBoost

Fit third stump


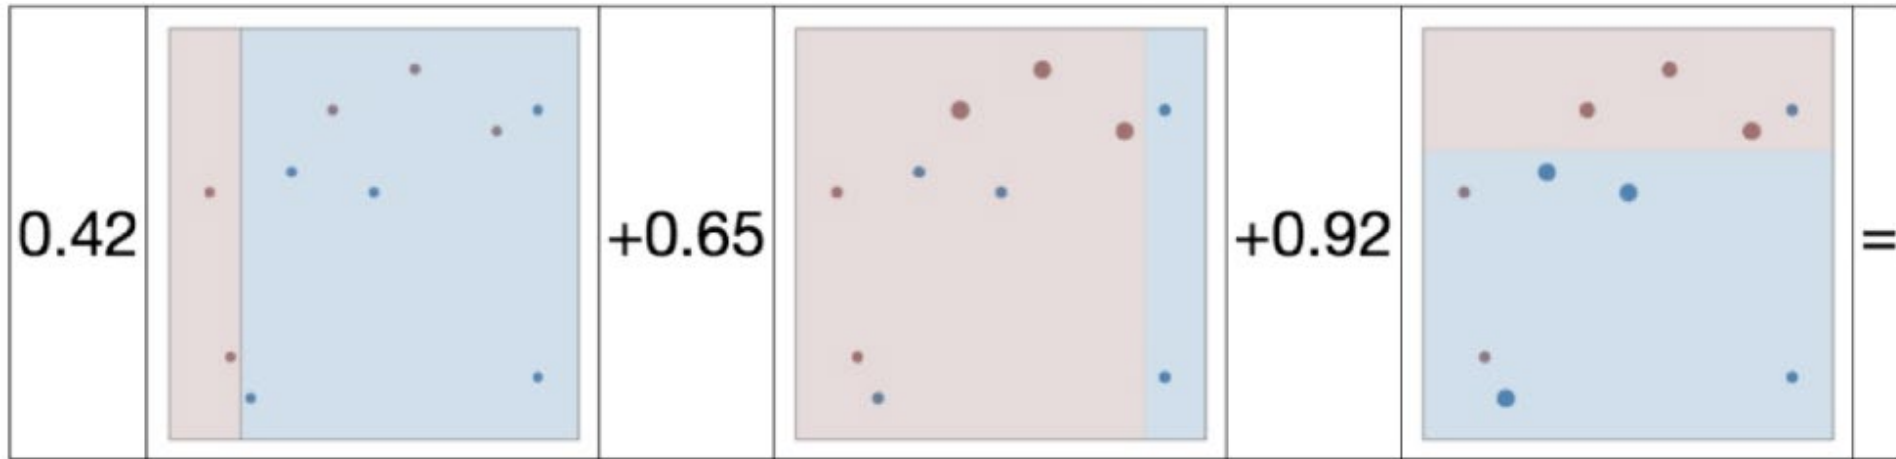


$k = 2$, $\epsilon = 0.14$, $\lambda = 0.92$



AdaBoost

The result:



Popular Boosted Tree Methods

- AdaBoost (Adaptive Boosting)
- GBM (Gradient Boosting Machine)
- XGBoost (Extreme Gradient Boosting)

Gradient Boosting

Initialize $F_0(\mathbf{x}) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$

For $k=1$ to K :

For $i=1,2,\dots,N$ compute

$$r_{ik} = - \left[\frac{\partial L(y_i, F_{k-1}(\mathbf{x}_i))}{\partial F_{k-1}(\mathbf{x}_i)} \right]$$

$$y_i - y_{pr}$$

$$\frac{1}{2}(y-f)^2 \quad y-f$$

Fit a regression tree to the targets r_{ik}

For the terminals $j=1,\dots,m$ compute

$$\gamma_{jk} = \arg \min_{\gamma} \sum_{\mathbf{x}_i \in R_{jk}} L(y_i, F_{k-1}(\mathbf{x}_i) + \gamma)$$

Update $F_k(\mathbf{x}) = F_{k-1}(\mathbf{x}) + \sum_{j=1}^m \gamma_{jk} I(\mathbf{x} \in R_{jk})$

Output $\hat{f}(\mathbf{x}) = F_K(\mathbf{x})$

Gradient Boosting

Setting	Loss Function	$-\partial L(y_i, f(x_i))/\partial f(x_i)$
Regression	$\frac{1}{2}[y_i - f(x_i)]^2$	$y_i - f(x_i)$
Regression	$ y_i - f(x_i) $	$\text{sign}[y_i - f(x_i)]$
Regression	Huber	$y_i - f(x_i)$ for $ y_i - f(x_i) \leq \delta_m$ $\delta_m \text{sign}[y_i - f(x_i)]$ for $ y_i - f(x_i) > \delta_m$ where $\delta_m = \alpha\text{th-quantile}\{ y_i - f(x_i) \}$
Classification	Deviance	$k\text{th component: } I(y_i = \mathcal{G}_k) - p_k(x_i)$

Popular Boosted Tree Methods

- AdaBoost (Adaptive Boosting)
- GBM (Gradient Boosting Machine)
- XGBoost (Extreme Gradient Boosting)

XGBoost

- Uses Newton tree boosting method
- Implements regularization helping reduce overfit (GB does not have)
- Implements parallel processing being much faster than GB

Another Variant: LightGB (very fast)

Python packages

```
class sklearn.ensemble.AdaBoostClassifier(base_estimator=None, n_estimators=50, learning_rate=1.0, algorithm='SAMME.R',  
random_state=None) \[source\]
```

```
class sklearn.ensemble.GradientBoostingClassifier(loss='deviance', learning_rate=0.1, n_estimators=100, subsample=1.0,  
criterion='friedman_mse', min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_depth=3,  
min_impurity_decrease=0.0, min_impurity_split=None, init=None, random_state=None, max_features=None, verbose=0,  
max_leaf_nodes=None, warm_start=False, presort='deprecated', validation_fraction=0.1, n_iter_no_change=None, tol=0.0001,  
ccp_alpha=0.0) \[source\]
```

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html>

https://xgboost.readthedocs.io/en/latest/python/python_intro.html