


Neural Networks (2)


Geena Kim



Perceptron




- Binary Threshold (Step function)




$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

Activation functions




- Binary Threshold (Step function)




$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

- Sigmoid



- Tanh



Perceptron

What can a perceptron do?

- Boolean tasks
- Update the weights whenever the perceptron output is wrong
- Proved convergence for linearly separable classes

Perceptron

Learning in perceptron

- Perceptron rule
- Delta rule (Gradient Descent)


Perceptron

Perceptron rule

$$\omega_j \leftarrow \omega_j - \alpha (\hat{y}_i - \underline{y_i}) X_{ij}$$

Handwritten annotations for the equation:

- Red arrow pointing to α with label "1"
- Red circle around ω_j in the update term
- Red circle around X_{ij} with label "feature column"
- Red circle around \hat{y}_i with label "sample"
- Red text "- 2" above the minus sign
- Red text "+ 2" below the minus sign
- Key: Red -1, Blue = +1



Perceptron

Perceptron algorithm

- Cycle through the training instances
- Only update W on misclassified instances
- If instance misclassified:
 - If instance is positive class (positive misclassified as negative)

$$W = W + X_i$$

- If instance is negative class (negative misclassified as positive)

$$W = W - X_i$$


$$w'_j \leftarrow w_j + \underbrace{1 (\hat{y}_i - y_i)}_{\hat{y}_i = 1, y_i = -1} X_{ij}$$

$\hat{y}_i = 1$
 $y_i = -1$
 $\hat{y}_i = 1$
 $y_i = -1$
 $y_i = +1$

$w_j =$
 $(w_j - 2)$
 w_j
 $w_j + 2$

Perceptron


Perceptron algorithm



Perceptron

Perceptron algorithm

ΔW




Updated weight vector

Misclassified *positive* instance, *add* it to W


Perceptron

Perceptron algorithm




Perceptron

Perceptron algorithm



Perceptron


Perceptron algorithm



Misclassified *negative* instance, *subtract* it from W

Perceptron

Perceptron algorithm



Perceptron

Delta rule (Gradient Descent)

$$\omega_j \leftarrow \omega_j - \alpha \frac{\partial \mathcal{L}}{\partial \omega_j}$$

$\mathcal{L}(\omega_j, x_{ji})$

$$MSE = \sum_i \frac{1}{2} (y_i - \hat{y}_i)^2$$
$$\frac{\partial}{\partial w} (y - (wx + b))^2$$


$$\mathcal{L} = \frac{1}{2} (\hat{y}_i - y_i)^2$$
$$- \alpha (\hat{y} - y) x$$

$$(y - \underbrace{(wx + b)}_{\hat{y}}) \cdot (-x)$$

$$\hat{y}_i = \sum_j \omega_j X_{ij}$$

$$\omega_j \leftarrow \omega_j - \alpha (\hat{y}_i - y_i) X_{ij}$$


No solution for XOR!
Not universal!



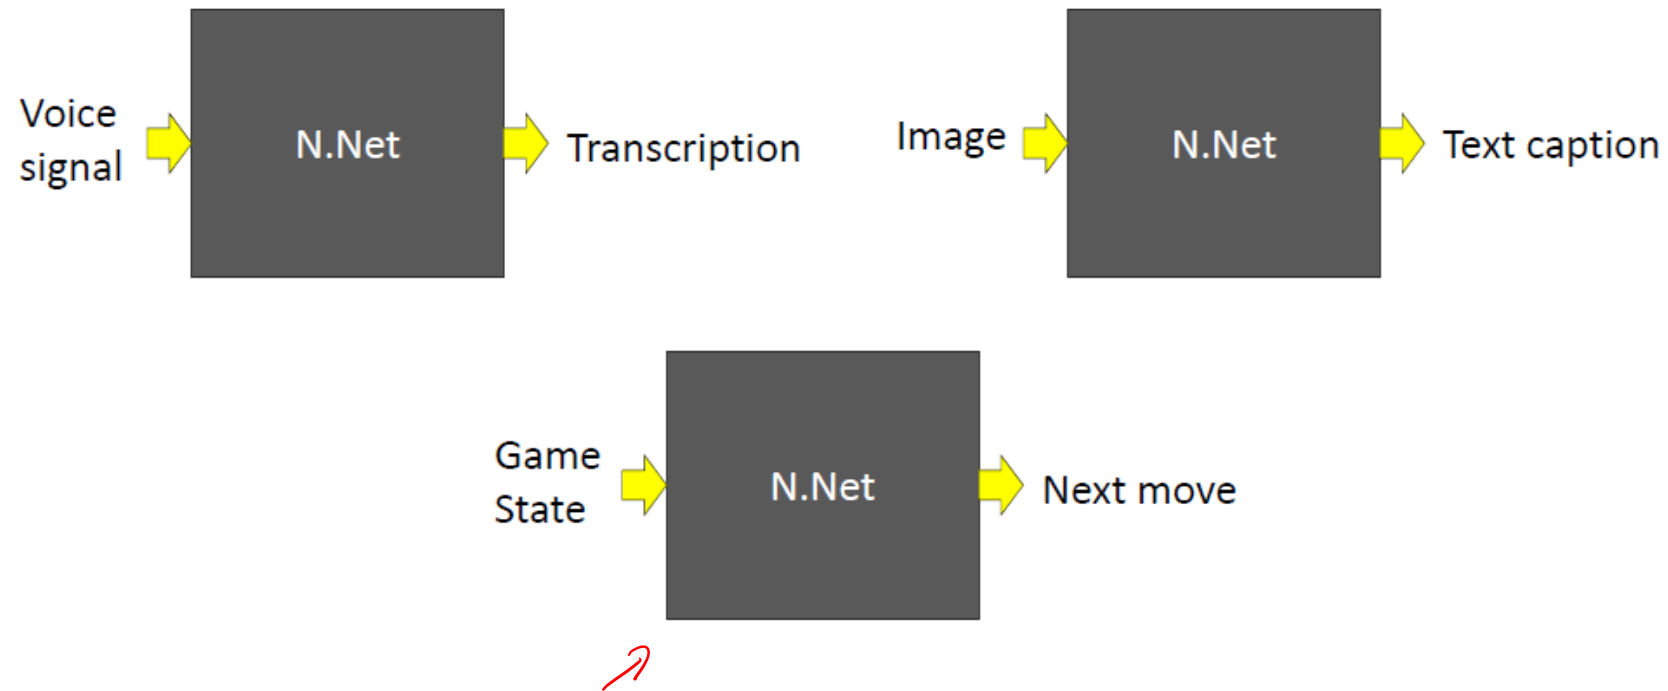
- Minsky and Papert, 1968

MLP


How do we handle linearly inseparable cases?



Neural networks are universal function approximator




We can make an arbitrary shape decision boundary




Can now be composed into
"networks" to compute arbitrary
classification "boundaries"

Boolean over real numbers




The network must fire if the input is in the coloured area

Boolean over real numbers




The network must fire if the input is in the coloured area

Boolean over real numbers





The network must fire if the input is in the coloured area

Boolean over real numbers




The network must fire if the input is in the coloured area

Boolean over real numbers




The network must fire if the input is in the coloured area

Boolean over real numbers




The network must fire if the input is in the coloured area

Neural networks can be used for regression tasks




Multi-layer Perceptron (Neural network)



Design Parameters

- Architecture
- Number of layers
- Number of neurons in a layer
- Activation functions




Effect of depth in network architecture

Theorem 5

A certain class of functions \mathcal{F} of n inputs can be represented using a deep network with $\mathcal{O}(n)$ units, whereas it would require $\mathcal{O}(2^{\sqrt{n}})$ units for a shallow network.

Theorem 6

For a certain class of functions \mathcal{G} of n inputs, the deep sum-product network with depth k can be represented with $\mathcal{O}(nk)$ units, whereas it would require $\mathcal{O}((n-1)^k)$ units for a shallow network.



Shallow vs. Deep Sum-Product Networks


Olivier Delalleau

Department of Computer Science and Operation Research
Université de Montréal
delallea@iro.umontreal.ca

Yoshua Bengio

Department of Computer Science and Operation Research
Université de Montréal
yoshua.bengio@umontreal.ca

Next time: How Neural Network Training Works



Next time: More on Weight Update in deep neural nets

$$W_{ij} \leftarrow W_{ij} - \alpha \frac{\partial \mathcal{L}}{\partial W_{ij}}$$

Weight Update Rule

Loss Function

Gradient (Chain Rule)

Back Propagation

Play with Neural network training!

You can change the model architecture, model hyper parameters, train hyperparameters. See what happens when you make each change.

<https://colab.research.google.com/drive/1lctSzyugG9YLJKhg7WAlRtGi0nElxy9y>